

Introduction to Cloud Computing (CS 524)

(Homework 3)

Prof. Igor Faynberg

Student Name: **Paras Garg**

Course Section: **CS 524-A**

Homework 3.1 –

Given the token bucket size, b bytes; token rate, r bytes/sec; and maximum output rate M bytes/sec. What is the maximum burst time T ?

Solution 3.1 –

The token bucket algorithm is based on an analogy of a fixed capacity bucket into which tokens, normally representing a unit of bytes or a single packet of predetermined size, are added at a fixed rate. When a packet is to be checked for conformance to the defined limits, the bucket is inspected to see if it contains sufficient tokens at that time. If so, the appropriate number of tokens (length of packet in bytes), are removed and the packet is passed (for transmission).

Steps to calculate burst time are as follows:

1. A token is added to the bucket every $\frac{1}{r}$ seconds.
2. It is given that the bucket size is b bytes means the bucket can hold at the most b tokens and if a token arrives when the bucket is full will be discarded.
3. If a packet of ' n bytes' arrives then ' n ' tokens will be removed from the bucket and the packet is sent to the network. And if fewer than ' n ' tokens are available then no tokens will be removed from the bucket.
4. Now, considering M is a maximum output rate and r is a token rate then T_{\max} (maximum burst time) will be:

$$T_{\max} = \begin{cases} b/(M - r) & \text{if } r < M \\ \infty & \text{otherwise} \end{cases}$$

5. In the above equation, T_{\max} is the time for which the rate M is fully utilized. Therefore, the maximum burst time is $T_{\max} = b/(M-r)$.

Example to calculate burst time:

Let bucket size, $b = 10$ bytes; token rate, $r = 1$ bytes/sec; and maximum output rate, $M = 2$ bytes/sec.

$$\begin{aligned} T &= b / (M - r) \\ T &= 10 / (2 - 1) \\ T &= 10 / 1 \\ T &= 10 \text{ sec} \end{aligned}$$

(Reference: https://en.wikipedia.org/wiki/Token_bucket)

Homework 3.2 –

Study the *AWS Direct Connect* service and answer the following questions:

- a. You own a company with a data center in Sapporo, Japan. Which company would you choose to connect this location to the Amazon service? Can you find out about pricing and QoS guarantees? (This may require a good deal of research. If you are unable to find the exact answers, describe what you have done to find them and what remains to be done.)
- b. As you have noticed, the *AWS Direct Connect* service description refers to the IEEE standard 802.1q. Read this standard (available at <http://standards.ieee.org/getieee802/download/802.1Q-2005.pdf>) and explain how a dedicated connection can be partitioned into multiple virtual interfaces so as to allow you to "use the same connection to access public resources such as objects stored in Amazon S3 using public IP address space, and private resources such as Amazon EC2 instances running within an Amazon Virtual Private Cloud (VPC) using private IP space."

Solution 3.2 –

AWS Direct Connect makes it easy to establish a dedicated network connection from premises to AWS. We can establish private connectivity between AWS and data center, office, co-location environment, which in many cases can reduce network costs, increase bandwidth throughput, and provide a more consistent network experience than Internet based connections.

- a. If I own a company with data center in Sapporo, Japan. I would choose to connect **Equinix, Inc** among other partners to connect this location to the Amazon Web Services. As with Equinix data centers
 - Offers many options for high-performance, private access to AWS Direct Connect.
 - Depending upon the data volume, AWS Direct connect customers can cut data transfer costs by two to ten times.

- Provides a flexible range of speeds (50, 100, 200, 300, 400, and 500 Megabits per second) with Virtual Connections (VCs) via Equinix Cloud Exchange.
- Provides cross connections either 1 or 10 Gigabits per second connections.
- Offers Amazon Direct Connect services covers more geographical locations than any other data center provider. AWS GovCloud is available in all U.S. AWS Direct Connect locations.
- Offers latest technology capabilities and has the facility to migrate to hybrid cloud computing.

QoS Guarantees

- **Reliability** – All Equinix IBX data centers are equipped with full UPS power, back-up systems and N+1 (or greater) redundancy, with a proven, industry-leading >99.99999% uptime record.
- **Power Density** – With robust heating, ventilation and air conditioning (HVAC) systems, Equinix IBX data centers exceed the requirements of even the most power-hungry deployments.
- **Security** – Each Equinix IBX data center utilizes an array of security equipment, techniques and procedures to control, monitor, and record access to the facility, including individual cages.
- **Recovery** – IBXflex™ Space provides operations centers and storage space when you need it; Equinix Smart Hands™ offers 24-hour access to qualified technical support—with Equinix, you can maintain your mission-critical operations and equipment under any circumstances.
- **High average uptime**—Our IBX data centers boast an industry-leading track record of >99.99999%.
- **Proven expertise**—We can help you configure and support your high-power density deployments.
- **The largest worldwide footprint**—We provide 14M+ ft² of data center space across 40 markets in 21 countries on 5 continents.
- **More than 1,400+ networks**—We've got you covered for interconnection and traffic exchange via direct cross connects, peering and cloud services.

AWS Direct Connect Pricing

Port Speed	Port-Hour Rate (All Regions except Japan)	Port-Hour Rate in Japan
50M*	\$0.03/hour	\$0.029/hour
100M*	\$0.06/hour	\$0.057/hour
200M*	\$0.12/hour	\$0.114/hour
300M*	\$0.18/hour	\$0.171/hour
400M*	\$0.24/hour	\$0.228/hour
500M*	\$0.30/hour	\$0.285/hour
1G	\$0.30/hour	\$0.285/hour
10G	\$2.25/hour	\$2.142/hour

(Reference: <https://aws.amazon.com/directconnect/partners/#apac>, <http://www.equinix.com/partners/amazon-web-services/>, <http://www.equinix.com/services/data-centers-colocation/>, <https://aws.amazon.com/directconnect/pricing/>)

- With AWS Direct Connect, you can establish 1 Gbps or 10 Gbps dedicated network connections between AWS and any of the AWS Direct Connect locations. A dedicated connection can be partitioned into multiple logical connections by using industry standard 802.1QVLANS. In this way, you can use the same connection to access public resources, such as objects stored in Amazon Simple Storage Service (Amazon S3) that use public IP address space, and private resources, such as Amazon EC2 instances that are running within a VPC using private IP space—all while maintaining network separation between the public and private environments. You can choose a partner from the AWS Partner Network (APN) to integrate the AWS Direct Connect endpoint in an AWS Direct Connect location with your remote networks.

Finally, you may combine all these different options in any combination that make the most sense for your business and security policies. For example, you could attach a VPC to your existing data centre with a virtual private gateway and set up an additional public subnet to connect to other AWS services that do not run within

the VPC, such as Amazon S3, Amazon Simple Queue Service (Amazon SQS) or Amazon Simple Notification Service (Amazon SNS).

AWS will allocate private IPs (/30) in the 169.x.x.x range for the BGP session and will advertise the VPC CIDR block over BGP. We can advertise the default route via BGP. A VPC VPN Connection creates encrypted network connectivity between your intranet and Amazon VPC over the Internet with the help of IPSec.

VPN Connections can be configured quickly, have low to modest bandwidth requirements, and can tolerate the inherent variability in Internet-based connectivity. AWS Direct Connect does not involve the Internet; instead, it uses dedicated, private network connections between your intranet and Amazon VPC.

(Reference: <http://docs.aws.amazon.com/directconnect/latest/UserGuide/dc-ug.pdf>,
<https://aws.amazon.com/directconnect/faqs/>)

Homework 3.3 –

Describe how the *AWS Direct Connect* service can be used with the *Amazon Virtual Private Cloud (VPC)*.

Solution 3.3 –

AWS Direct Connect links your internal network to an AWS Direct Connect location over a standard 1 gigabit or 10 gigabit Ethernet fiber-optic cable. One end of the cable is connected to your router, the other to an AWS Direct Connect router.

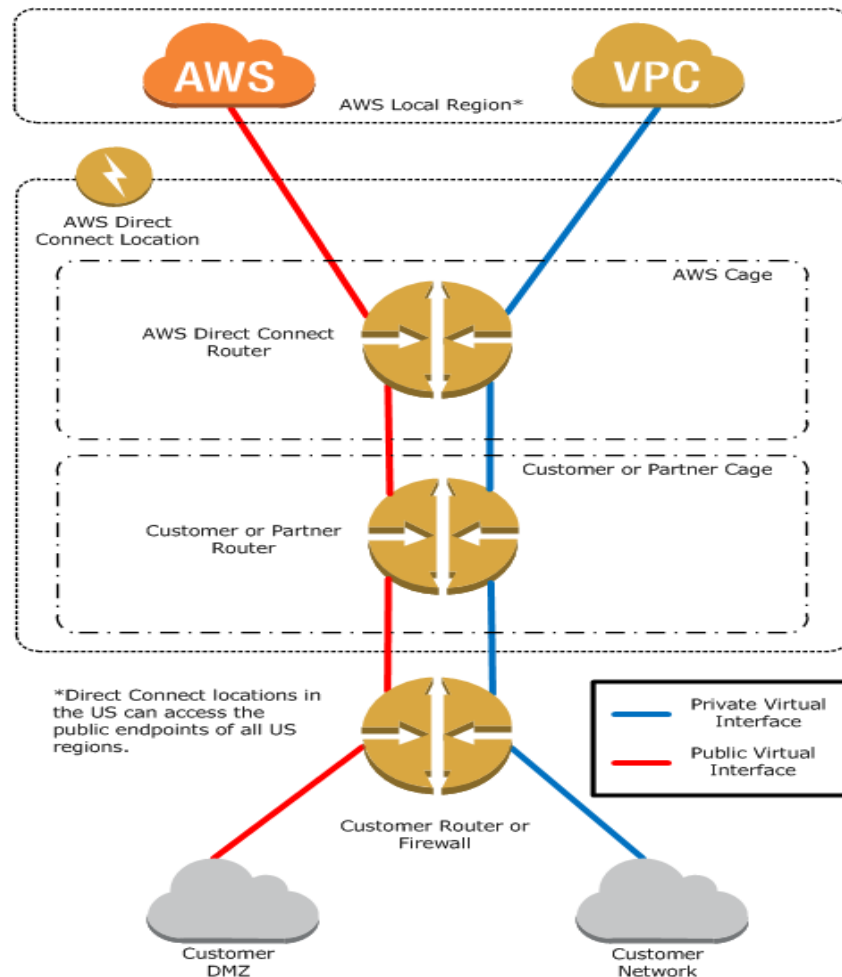
You can create virtual interfaces directly to the AWS cloud (for example, to Amazon EC2 and Amazon S3) and to Amazon VPC, bypassing Internet service providers in your network path. When you create a private virtual interface to a VPC, you will need a private virtual interface for each VPC you want to connect. This connection requires the use of Border Gateway Protocol (BGP). You need the following information to complete the connection:

- A public or private ASN. If you are using a public ASN you must own it else if you are using a private ASN, it must be in the 65000 range.
- A new unused VLAN tag that you select
- The VPC Virtual Private Gateway (VGW) id (AWS will allocate private IPs (/30) in the 169.x.x.x range for the BGP session and will advertise the VPC CIDR block over BGP. You can advertise the default route via BGP).

The Virtual Private Gateway to connect:

- Verify that the VLAN is not already in use on this connection.
- Open the AWS Direct Connect console.
- In the connection pane, select the connection to use, and then click *Create Virtual Interface*.
- In the *Create a Virtual Interface* pane, select *Private*.
- Under *Define Your New Private Virtual Interface*
 - Enter a name for the virtual interface in the *Interface Name* field.
 - In the *Interface Owner*, select the *My AWS Account* option if the virtual interface is for your AWS account ID.
 - Select the virtual gateway to connect to, in the *VGW* list.
 - In the *VLAN #* field, enter the ID number for your virtual local area network (VLAN); for example, a number between 1 and 4094.
 - To have AWS generate your router IP address and Amazon IP address, select *Auto-generate peer IPs*.
 - To specify these IP addresses yourself, clear the *Auto-Generated peer IPs* check box, and then in the *Your router peer IP* field, enter the destination IPv4 CIDR address that Amazon should send traffic to. In the *Amazon router peer IP* field, enter the IPv4 CIDR address you will use to Amazon Web Services.
 - In the *BGP ASN* field, enter the Border Gateway Protocol (BGP) Autonomous System Number (ASN) of your gateway; for example, a number between 1 and 65534.
 - Select *Auto-generate BGP key* check box to have AWS generate one.
- To provide your own BGP key, clear the **Auto-generate BGP key** check box, and then in the **BGP Authorization Key** field, enter your BGP MD5 key.
- Then download your router configuration and configure the router.

An AWS Direct Connect location provides access to Amazon Web Services in the region it is associated with, as well as access to other US regions. For example, you can provision a single connection to any AWS Direct Connect location in the US and use it to access public AWS services in all US Regions and AWS GovCloud (US). The following diagram shows how AWS Direct Connect interfaces with your network.



(Reference: <http://docs.aws.amazon.com/directconnect/latest/UserGuide/Welcome.html>)

Homework 3.4 –

Note that *Amazon VPC* provides NAT.

- Explain why you would want to use *NAT* for a *virtual private subnet with the Amazon Direct Connect* service. Do you see any cases where you would *not* want to use it?
- What is the maximum number of connections a single NAT box can maintain? (You need to check the specifications of the three existing transport-layer protocols on the Internet: TCP, UDP, and SCTP, and also keep in mind that the first 4,096 ports have been reserved.)

Solution 3.4 –

Network Address Translation (NAT) is a method of remapping one IP address space into another by modifying network address information in Internet Protocol (IP) datagram packet headers while they are in transit across a traffic routing device. For example, router, which act as an agent between a local (or private network) and the internet (or public network), which represent an unique IP address to public network on the behalf of entire group of computers on private network.

- We use NAT for virtual private subnet with the Amazon Direct Connect Services to enable instances in a private subnet to connect to the Internet (for example, for software updates) or other AWS services, but prevent the Internet from initiating connections with instances. A NAT forwards traffic from the instances in the private

subnet to the Internet or other AWS services, and then sends the response back to the instances. When traffic goes to the Internet, the source IP address is replaced with the NAT device's address and similarly, when response traffic goes to those instances, the NAT device translates the address back to those instances' private IP addresses.

The cases where the uses of NAT are bad, when a user is working with instances that require the use of static public IP address and when there is no Internet gateway to enable communication over the Internet as this scenario includes a virtual private cloud (VPC) with a single private subnet, and a virtual private gateway to enable communication with own network over an IPsec VPN tunnel.

(Reference: https://en.wikipedia.org/wiki/Network_address_translation, <http://computer.howstuffworks.com/nat.htm>, <http://docs.aws.amazon.com/AmazonVPC/latest/UserGuide/vpc-nat.html>)

- b. The maximum number of connections that a single NAT box can maintain is 2^{16} ie 65,536. But it is known that the first 4,096 ports are reserved, therefore, the effective number of maximum connections that can be used are 65,536-4096 ie **61440**.

(Reference: Cloud Computing: Business Trends and Technologies)

Homework 3.5 –

Read RFC 1930 (<http://www.ietf.org/rfc/rfc1930.txt>) and answer the following questions:

- a. To use AWS Direct Connect with *Amazon VPC*, the Border Gateway Protocol is required. Why?
- b. Can you use your own *ASN* to connect to VPC?
- c. Which RIR would you go to when you need to establish an ASN for your data center in Sapporo, Japan?

Solution 3.5 –

- a. To use AWS Direct Connect with Amazon VPC required the use of the Border Gateway Protocol (BGP) with an Autonomous System Number (ASN) and IP Prefixes.

Border Gateway Protocol is a standardized exterior gateway protocol designed to exchange routing and reachability information among autonomous systems (AS) on the internet. BGP is used to communicate between two routing domains (to exchange routes from VPC to private network). BGP enables sending router decides on the shortest path to the destination based on the routing table lookup which was previously obtained from a "neighbor" and subsequently updated.

(References: https://en.wikipedia.org/wiki/Border_Gateway_Protocol, <https://aws.amazon.com/directconnect/faqs/>)

- b. An Autonomous System Number (ASN) is a collection of connected Internet Protocol (IP) routing prefixes under the control of one or more network operators. ASN are used to identify networks that present common, clearly defined routing policy to the Internet.

Yes, we can use our own ASN to connect to VPC

AWS Direct Connect requires an ASN to create a public or private virtual interface. We may use a public ASN which we own or we can pick any private ASN between 64512 to 65534.

(References: [https://en.wikipedia.org/wiki/Autonomous_system_\(Internet\)](https://en.wikipedia.org/wiki/Autonomous_system_(Internet)), <https://aws.amazon.com/directconnect/faqs/>)

- c. Regional Internet Registry (RIR) is an organization that manages the allocation and registration of Internet number resources within a particular region of the world. Internet number resources include IP addresses and autonomous system (AS) numbers.

To establish an ASN for my data center in Sapporo, Japan, I would go to Asia Pacific Network Information Centre (APNIC). APNIC provides number resource allocation and registration services that support the global operation of the Internet. It is a non-for-profit, membership-based organization whose members include Internet Service Providers, National Internet Registries, and similar organizations.

The main functions of APNIC are:

- allocating IPv4 and IPv6 address space, and Autonomous System Numbers,
- maintaining the public Whois Database for the Asia Pacific region,
- reverse DNS delegations,
- representing the interests of the Asia Pacific Internet community on the global stage.

(References: https://en.wikipedia.org/wiki/Regional_Internet_registry, https://en.wikipedia.org/wiki/Asia-Pacific_Network_Information_Centre, <https://www.apnic.net/>)